

Business Intelligence

Lecture 5: **Text and Web Mining**

Dr Tamal Ghosh

Learning Objectives

- Describe text mining and understand the need for text mining
- Differentiate between text mining, Web mining and data mining
- Understand the different application areas for text mining
- Know the process of carrying out a text mining project
- Understand the different methods to introduce structure to text-based data

Learning Objectives

- Describe Web mining, its objectives, and its benefits
- Understand the three different branches of Web mining
 - Web content mining
 - Web structure mining
 - Web usage mining
- Understand the applications of these three mining paradigms

Opening Vignette...

“Mining Text For Security And Counterterrorism”

- What is MITRE?
- Problem description
- Proposed solution
- Results
- Answer & discuss the case questions

Opening Vignette: Mining Text For Security...

Cluster 1

(L) Kampala
(L) Uganda
(P) Yoweri Museveni
(L) Sudan
(L) Khartoum
(L) Southern Sudan

Cluster 2

(P) Timothy McVeigh
(L) Oklahoma City
(P) Terry Nichols

Cluster 3

(E) election
(P) Norodom Ranariddh
(P) Norodom Sihanouk
(L) Bangkok
(L) Cambodia
(L) Phnom Penh
(L) Thailand
(P) Hun Sen
(O) Khmer Rouge
(P) Pol Pot

Text Mining Concepts

- 85-90 percent of all corporate data is in some kind of unstructured form (e.g., text).
- Unstructured corporate data is doubling in size every 18 months.
- Tapping into these information sources is not an option, but a need to stay competitive.
- Answer: text mining
 - A semi-automated process of extracting knowledge from unstructured data sources
 - a.k.a. text data mining or knowledge discovery in textual databases

Data Mining versus Text Mining

- Both seek novel and useful patterns
- Both are semi-automated processes
- Difference is the nature of the data:
 - Structured versus unstructured data
 - **Structured data:** databases
 - **Unstructured data:** Word documents, PDF files, text excerpts, XML files, and so on
- Text mining – first, impose structure to the data, then mine the structured data

Text Mining Concepts

- Benefits of text mining are obvious especially in text-rich data environments
 - e.g., law (court orders), academic research (research articles), finance (quarterly reports), medicine (discharge summaries), biology (molecular interactions), technology (patent files), marketing (customer comments), etc.
- Electronic communication records (e.g., Email)
 - Spam filtering
 - Email prioritization and categorization
 - Automatic response generation

Text Mining Application Area

- Information extraction
- Topic tracking
- Summarization
- Categorization
- Clustering
- Concept linking
- Question answering

Text Mining Terminology

- Unstructured or semistructured data
- Corpus (and corpora)
- Terms
- Concepts
- Stemming
- Stop words (and include words)
- Synonyms (and polysemes)
- Tokenizing

Text Mining Terminology

- Term dictionary
- Word frequency
- Part-of-speech tagging
- Morphology
- Term-by-document matrix
 - Occurrence matrix
- Singular value decomposition
 - Latent semantic indexing

Text Mining for Patent Analysis (see Applications Case 7.2)

- What is a patent?
 - “exclusive rights granted by a country to an inventor for a limited period of time in exchange for a disclosure of an invention”
- How do we do patent analysis (PA)?
- Why do we need to do PA?
 - What are the benefits?
 - What are the challenges?
- How does text mining help in PA?

Natural Language Processing (NLP)

- Structuring a collection of text
 - Old approach: bag-of-words
 - New approach: natural language processing
- NLP is
 - a very important concept in text mining.
 - a subfield of artificial intelligence and computational linguistics.
 - the study of "understanding" the natural human language.
- Syntax versus semantics based text mining

Natural Language Processing (NLP)

- What is “Understanding” ?
 - Human understands, what about computers?
 - Natural language is vague, context driven
 - True understanding requires extensive knowledge of a topic
- Can/will computers ever understand natural language the same/accurate way we do?

Natural Language Processing (NLP)

- Challenges in NLP
 - Part-of-speech tagging
 - Text segmentation
 - Word sense disambiguation
 - Syntax ambiguity
 - Imperfect or irregular input
 - Speech acts
- Dream of AI community
 - to have algorithms that are capable of automatically reading and obtaining knowledge from text

Natural Language Processing (NLP)

- WordNet
 - A laboriously hand-coded database of English words, their definitions, sets of synonyms, and various semantic relations between synonym sets
 - A major resource for NLP
 - Needs automation to be completed
- Sentiment Analysis
 - A technique used to detect favorable and unfavorable opinions toward specific products and services
 - See Application Case 7.3 for a CRM application

NLP Task Categories

- Information retrieval
- Information extraction
- Named-entity recognition
- Question answering
- Automatic summarization
- Natural language generation & understanding
- Machine translation
- Foreign language reading & writing
- Speech recognition
- Text proofing
- Optical character recognition

Text Mining Applications


- Marketing applications
 - Enables better CRM
- Security applications
 - ECHELON, OASIS
 - Deception detection
 - example coming up
- Medicine and biology
 - Literature-based gene identification
 - example coming up
- Academic applications
 - Research stream analysis - example coming up

Text Mining Applications

- Application Case 7.4: Mining for Lies
- Deception detection
 - A difficult problem
 - If detection is limited to only text, then the problem is even more difficult
- The study
 - analyzed text based testimonies of persons of interest at military bases
 - used only text-based features (cues)

Text Mining Applications

- Application Case 7.4: Mining for Lies



Text Mining Applications

- Application Case 7.4: Mining for Lies

Category	Example Cues
Quantity	Verb count, noun-phrase count, ...
Complexity	Avg. no of clauses, sentence length, ...
Uncertainty	Modifiers, modal verbs, ...
Nonimmediacy	Passive voice, objectification, ...
Expressivity	Emotiveness
Diversity	Lexical diversity, redundancy, ...
Informality	Typographical error ratio
Specificity	Spatiotemporal, perceptual information ...
Affect	Positive affect, negative affect, etc.


Text Mining Applications

- Application Case 7.4: Mining for Lies
 - 371 usable statements are generated
 - 31 features are used
 - Different feature selection methods used
 - 10-fold cross validation is used
 - Results (overall % accuracy)

• Logistic regression	67.28
• Decision trees	71.60
• Neural networks	73.46


Text Mining Applications

(gene/protein interaction identification)




Text Mining Process

Context diagram for the text mining process



Text Mining Process



The three-step text mining process

Text Mining Process

- **Step 1:** Establish the corpus
 - Collect all relevant unstructured data (e.g., textual documents, XML files, emails, Web pages, short notes, voice recordings...)
 - Digitize, standardize the collection (e.g., all in ASCII text files)
 - Place the collection in a common place (e.g., in a flat file, or in a directory as separate files)

Text Mining Process

- **Step 2:** Create the Term-by-Document Matrix

Terms Documents	investment risk	project management	software engineering	development	SAP	...
Document 1	1			1		
Document 2		1				
Document 3			3		1	
Document 4		1				
Document 5			2	1		
Document 6	1			1		
...						

Text Mining Process

- **Step 2:** Create the Term-by-Document Matrix (TDM)
 - Should all terms be included?
 - Stop words, include words
 - Synonyms, homonyms
 - Stemming
 - What is the best representation of the indices (values in cells)?
 - Row counts; binary frequencies; log frequencies;
 - Inverse document frequency

Text Mining Process

- **Step 2:** Create the Term-by-Document Matrix (TDM)
 - TDM is a sparse matrix. How can we reduce the dimensionality of the TDM?
 - Manual – a domain expert goes through it
 - Eliminate terms with very few occurrences in very few documents (?)
 - Transform the matrix using singular value decomposition (SVD)
 - SVD is similar to principle component analysis

Text Mining Process

- **Step 2:** Extract patterns/knowledge
 - Classification (text categorization)
 - Clustering (natural groupings of text)
 - Improve search recall
 - Improve search precision
 - Scatter/gather
 - Query-specific clustering
 - Association
 - Trend Analysis (...)

Text Mining Application

(research trend identification in literature)

- Mining the published IS literature
 - MIS Quarterly (MISQ)
 - Journal of MIS (JMIS)
 - Information Systems Research (ISR)
- Covers 12-year period (1994-2005)
- 901 papers are included in the study
- Only the paper abstracts are used
- 9 clusters are generated for further analysis


Text Mining Application

(research trend identification in literature)

Journal	Year	Author(s)	Title	Vol/No	Pages	Keywords	Abstract
MISQ	2005	A. Malhotra, S. Gosain and O. A. El Sawy	Absorptive capacity configurations in supply chains: Gearing for partner- enabled market knowledge creation	29/1	145-187	knowledge management supply chain absorptive capacity interorganizational information systems configuration approaches	The need for continual value innovation is driving supply chains to evolve from a pure transactional focus to leveraging interorganizational partner ships for sharing
ISR	1999	D. Robey and M. C. Boudreau	Accounting for the contradictory organizational consequences of information technology: Theoretical directions and methodological implications	2-Oct	167-185	organizational transformation impacts of technology organization theory research methodology intraorganizational power electronic communication mis implementation culture systems	Although much contemporary thought considers advanced information technologies as either determinants or enablers of radical organizational change, empirical studies have revealed inconsistent findings to support the deterministic logic implicit in such arguments. This paper reviews the contradictory
JMIS	2001	R. Aron and E. K. Clemons	Achieving the optimal balance between investment in quality and investment in self- promotion for information products	18/2	65-88	information products internet advertising product positioning signaling signaling games	When producers of goods (or services) are confronted by a situation in which their offerings no longer perfectly match consumer preferences, they must determine the extent to which the advertised features of
...	...	...	...	...	...	...	...


Text Mining Application

(research trend identification in literature)



Text Mining Application

(research trend identification in literature)



Text Mining Tools


- Commercial Software Tools
 - SPSS PASW Text Miner
 - SAS Enterprise Miner
 - Statistica Data Miner
 - ClearForest
- Free Software Tools
 - RapidMiner
 - GATE
 - Spy-EM

Web Mining Overview

- Web is the largest repository of data
- Data is in HTML, XML, text format
- Challenges (of processing Web data)
 - The Web is too big for effective data mining
 - The Web is too complex
 - The Web is too dynamic
 - The Web is not specific to a domain
 - The Web has everything
- Opportunities and challenges are great!

Web Mining

- Web mining (or Web data mining) is the process of discovering intrinsic relationships from Web data (textual, linkage, or usage)



Web Content/Structure Mining

- Mining of the textual content on the Web
- Data collection via Web crawlers
- Web pages include hyperlinks
 - Authoritative pages
 - Hubs
 - hyperlink-induced topic search (HITS) alg.

Web Usage Mining


- Extraction of information from data generated through Web page visits and transactions
 - data stored in server access logs, referrer logs, agent logs, and client-side cookies
 - user characteristics and usage profiles
 - metadata, such as page attributes, content attributes, and usage data
- Clickstream data
- Clickstream analysis

Web Usage Mining

- Web usage mining applications
 - Determine the lifetime value of clients
 - Design cross-marketing strategies across products.
 - Evaluate promotional campaigns
 - Target electronic ads and coupons at user groups based on user access patterns
 - Predict user behavior based on previously learned rules and users' profiles
 - Present dynamic information to users based on their interests and profiles


Web Usage Mining

(clickstream analysis)



Web Mining Success Stories

- Amazon.com, Ask.com, Scholastic.com, etc.
- Website Optimization Ecosystem



Web Mining Tools

Product Name	URL
Angoss Knowledge WebMiner	angoss.com
ClickTracks	clicktracks.com
LiveStats from DeepMetrix	deepmetrix.com
Megaputer WebAnalyst	megaputer.com
MicroStrategy Web Traffic Analysis	microstrategy.com
SAS Web Analytics	sas.com
SPSS Web Mining for Clementine	spss.com
WebTrends	webtrends.com
XML Miner	scientio.com

End of the Chapter

- Questions, comments