NATURAL LANGUAGE PROCESSING

A foundational on bridging the gap between humans and machines

Dr Tamal Ghosh

MASTERING THE CORE LEARNING OBJECTIVES

Defining the scope and goals for today's introductory session

Defining NLP

Understand the field as an intersection of AI and linguistics.

Stating Importance

Identify why NLP is the backbone of modern search engines and AI assistants.

Major Applications

Catalog the diverse ways NLP solves real-world problems in business and speech.

Processing Pipeline

Map the journey of raw text through various preprocessing stages.

Linguistic Challenges

Recognize the inherent difficulties and ambiguities in human communication.

Shift in Paradigm

Differentiate between traditional programming logic and NLP learning models.

NATURAL LANGUAGE REMAINS DEEPLY AMBIGUOUS

Why human communication defies the rigid logic of traditional code

Feature	Natural Language	Programming Language
Clarity	Frequently ambiguous/context-heavy	Explicit and strictly defined
Evolution	Continuously changing with culture	Fixed syntax versions
Grammar	Rich, complex, and flexible	Strict, unforgiving syntax rules
Meaning	Depends on pragmatics and intent	Literal, deterministic execution

DEFINING NATURAL LANGUAGE PROCESSING

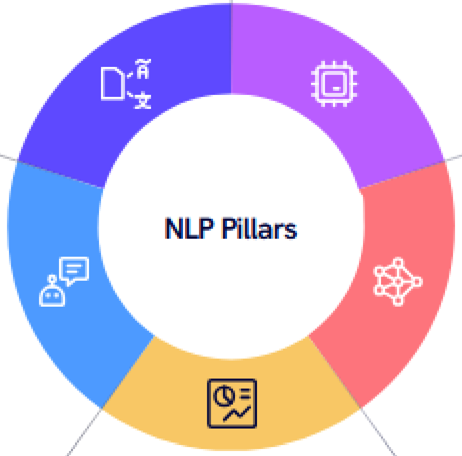

Linguistics
Provides structural/semantic rules.

Comp Science
Offers computational/algorithmic power.

Goal
Natural human-machine communication.

ML Systems
Systems learn from massive datasets.

Statistics
Probabilistic models for uncertainty.

NLP POWERS OUR DAILY DIGITAL INTERACTIONS

From global search to personal productivity tools

Information Retrieval

Google Search and Bing use NLP to understand intent rather than just keywords.

Generative AI

ChatGPT, Gemini, and Claude represent the current peak of conversational NLP.

Virtual Assistants

Siri and Alexa convert speech to text and then intent to action.

Security

Gmail's spam detection identifies malicious patterns in message content.

Accessibility

YouTube's automatic captions and real-time translation tools.

APPLICATION OVERVIEW

TAXONOMY OF REAL-WORLD NLP APPLICATIONS

Categorizing how NLP solves text, speech, and business challenges

Text-Based Tasks

Includes sentiment analysis for brand monitoring, machine translation for global communication, and document classification.

Speech-Based Tasks

Focuses on voice assistants, automated call center routing, and highly accurate speech-to-text transcription.

Business Logic

Utilizes NLP for customer support chatbots, legal document analysis, and extracting data from medical records.

Operational Tasking

Automated text summarization helps professionals digest long reports, while question-answering systems provide instant internal support.

CATALOGING COMMON ESSENTIAL NLP TASKS

The specific functions that form the building blocks of AI systems

Task	Description	Real-World Example
Translation	Converting text between languages	English to Bengali
Classification	Sorting text into predefined buckets	Identifying Spam vs. Not Spam
Sentiment Analysis	Determining emotional tone	Rating reviews as Positive/Negative
Summarization	Condensing long text into key points	Generating news bullet points
Named Entity Recognition	Identifying specific proper nouns	Extracting Person, Place, or Org
Text Generation	Creating new, coherent content	AI writing assistants like ChatGPT

MOVING FROM RULES TO LEARNED PATTERNS

The fundamental shift in how we program computers to understand language

Traditional Programming

Input + Rules = Output. Manually defining every rule fails due to language exceptions.

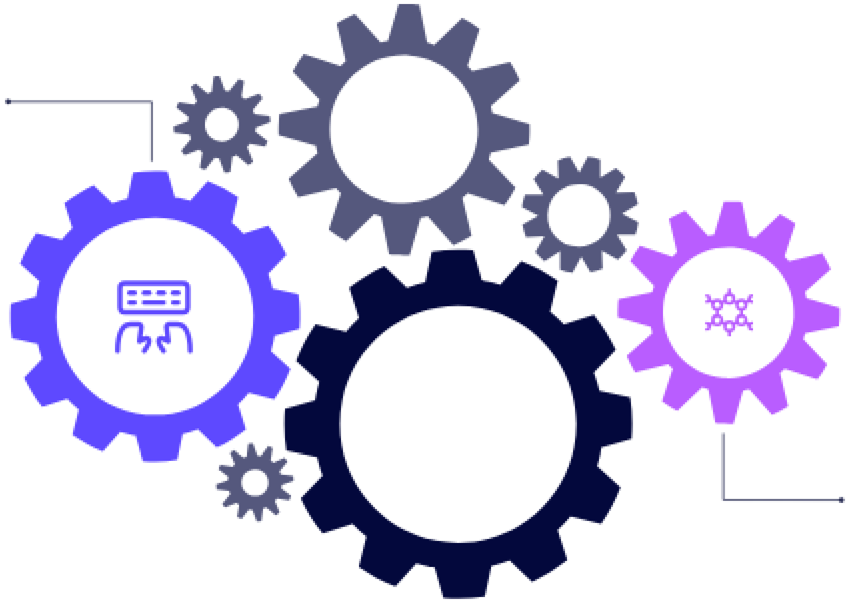

Artificial Intelligence

Input + Output = Learned Rules. Learning patterns and associations allows scalability.

presentations.ai

GREATEST CHALLENGE

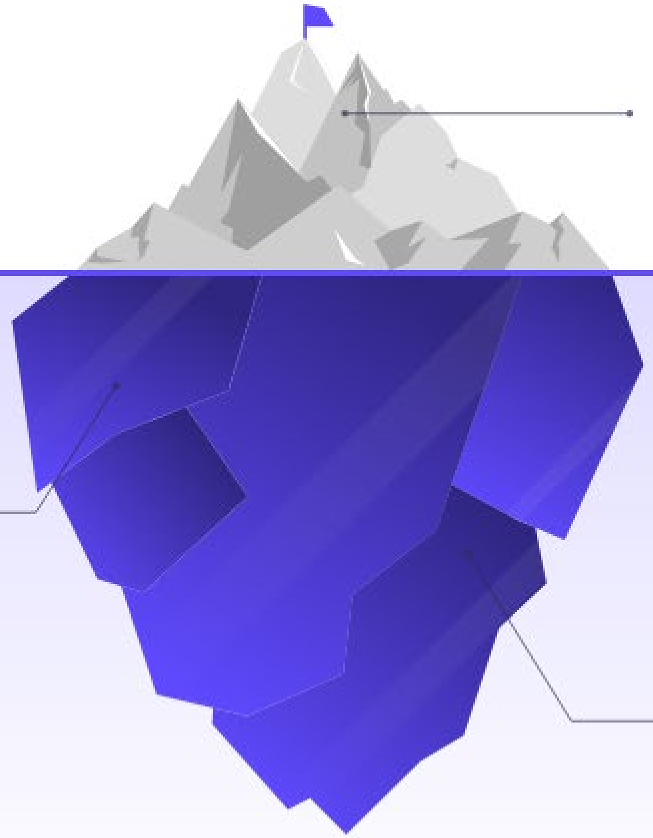

Empty Vessels
Words like 'bank' are ambiguous, meaning institution or river side.

Literal Data
Computers process only numeric values, identifying 'blue' as #0000FF.

Contextual Meaning
Humans interpret 'blue' as sadness, requiring probabilistic models.

LEXICAL AND SYNTACTIC AMBIGUITY BREAKDOWN

Structural hurdles that confuse even the best algorithms

01 **Lexical Ambiguity**

Occurs when a single word has multiple meanings. For example, in the sentence "I went to the bank," the system must determine if it refers to a financial bank or a river bank.

02 **Syntactic Ambiguity**

Occurs when the structure of a sentence allows for multiple interpretations. "Flying planes can be dangerous" could mean the act of piloting is dangerous, or that the planes are dangerous.

03 **Disambiguation**

Modern NLP uses context windows—looking at surrounding words—to calculate which meaning is statistically most likely.

SEMANTIC AND PRAGMATIC COMMUNICATION LAYERS

Understanding what is said versus what is actually meant

01 **Semantic Ambiguity**

Arises when the meaning of a whole sentence is unclear despite clear words. "Visiting relatives can be boring" leaves the machine wondering if the act of visiting is boring or if the people visiting are boring.

02 **Pragmatic Ambiguity**

The most difficult layer, where literal meaning differs from intended action. If someone says "It's cold here," a human knows to close the window, while a machine might just record the temperature (literal meaning).

03 **Contextual Awareness**

Solving these requires moving beyond individual words to understanding social cues and situational intent.

VISUALIZING THE STANDARD NLP PIPELINE

The linear journey from raw text to intelligent prediction

1

Segmentation

Breaking text into individual sentences.

2

Tokenization

Splitting sentences into individual words or "tokens."

3

Stop-word Removal

Filtering out common words (the, is, at) that carry little unique meaning.

4

Stemming/ Lemmatization

Reducing words to their root form (e.g., "learning" becomes "learn").

5

Feature Extraction

Converting these roots into numeric vectors.

6

ML Model

Applying algorithms to predict intent or generate a response.

PRACTICAL WALKTHROUGH OF TEXT PREPROCESSING

Transforming a live sentence into machine-readable data

Tokenization
Breaks the sentence into: The, students, are, learning, Natural, Language, Processing.

Stop-word Removal
Removes "The" and "are," leaving: students, learning, Natural, Language, Processing.

Stemming
Reduces words to their core: student, learn, natural, language, process.

NLU VS NLG: TWO SIDES OF THE SAME COIN

Dividing NLP into understanding and generation components

Natural Language Understanding

Focuses on the input side. Involves intent detection, sentiment analysis, and extracting structured information.

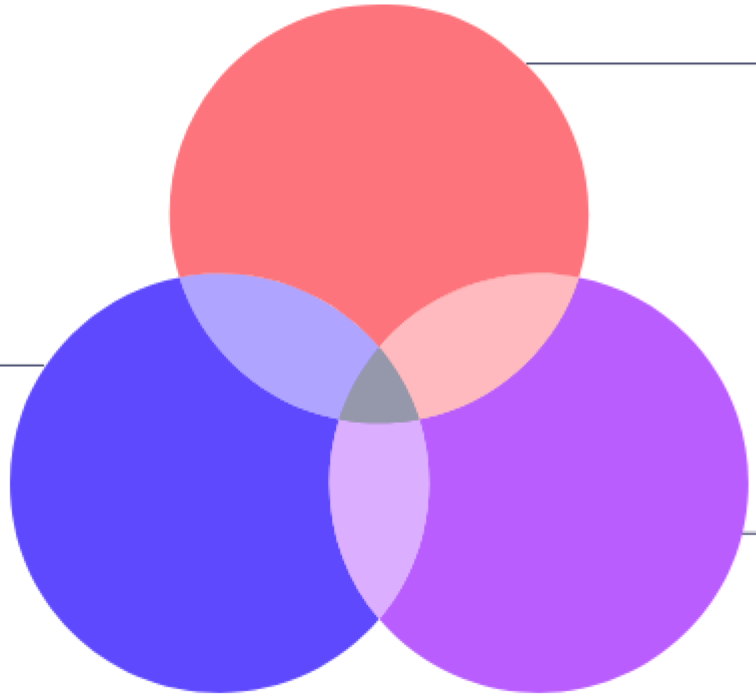

Natural Language Generation

Focuses on the output side. Process of producing human-like text from internal data representations.

The Conversational Loop

Modern agents like ChatGPT combine both: NLU parses the prompt and NLG crafts the answer.

THE DECADES-LONG EVOLUTION OF NLP

From rigid rules to massive billion-parameter models

1950s

Rule-Based

Initial attempts using strict linguistic rules and early machine translation.

1990s

Statistical NLP

Using probability to predict the next word or translate phrases.

2010

Machine Learning

The rise of neural networks and deep learning for text classification.

2018

Transformers (BERT)

A breakthrough architecture allowing models to understand long-range context in sentences.

2022+

Large Language Models (LLMs)

Massive models like GPT-4 and Llama that show emergent intelligence across thousands of tasks.

RULE-BASED VS MODERN MACHINE LEARNING

Why data-driven approaches have won the architectural war

Rule-Based Systems

Rely on human-written grammar rules; precise but difficult to maintain and scale to the trillions of possible sentence combinations.

VS

Machine Learning Systems

Learn directly from data; offer superior accuracy and handle complex, slang-filled, or grammatically incorrect language.

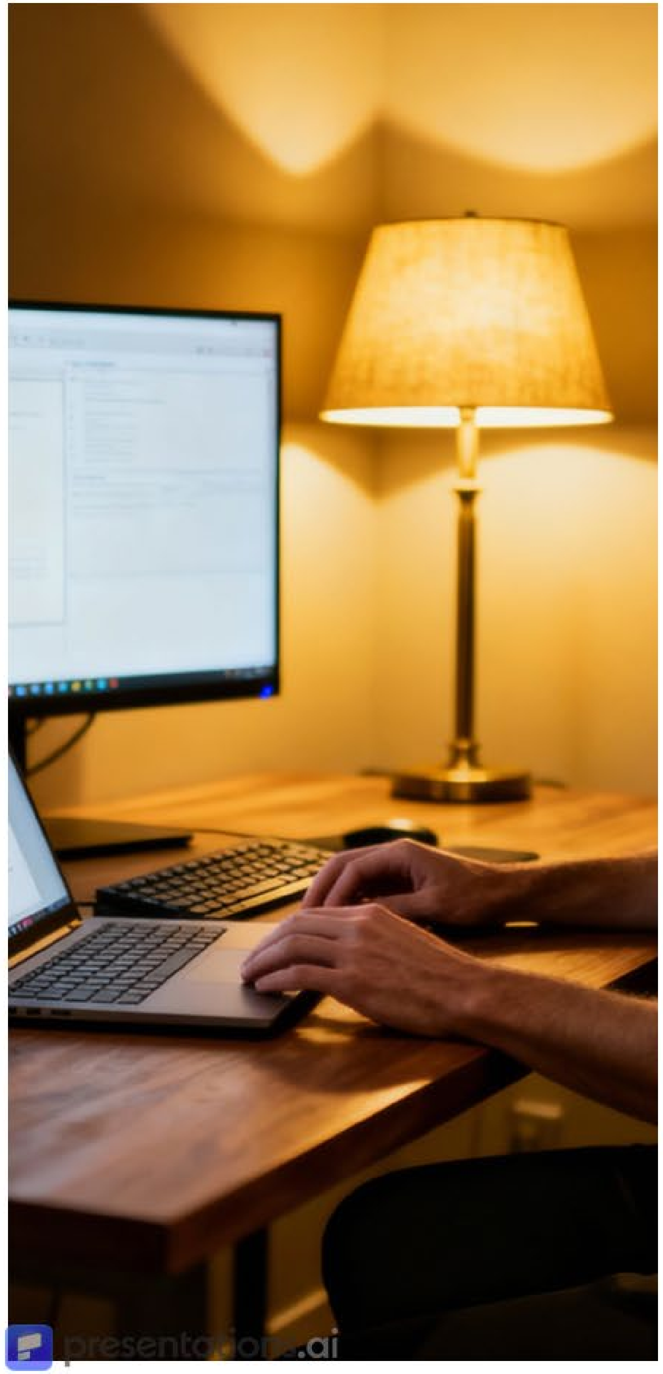

NAVIGATING THE PYTHON NLP ECOSYSTEM

Essential libraries for research and production

- **NLTK (Natural Language)**
  The "classic" library, excellent for learning and basic preprocessing like tokenization.
- **spaCy**
  Built for production; incredibly fast with advanced features like Named Entity Recognition out of the box.
- **Hugging Face Transformers**
  The current industry standard for accessing pre-trained LLMs and state-of-the-art models.
- **Gensim**
  Specialized in topic modeling and document similarity.
- **Scikit-learn**
  Useful for the machine learning layer of the pipeline, such as classification.

CASE STUDY: GMAIL SPAM DETECTION

A real-world breakdown of the NLP pipeline in action

Input

The system receives a message: "Congratulations! You have won ₹10,00,000. Click here."

Preprocessing

The text is tokenized, stop-words removed. Words like "won" and "click" are stemmed to their roots.

Feature Extraction

The system notices high-frequency keywords associated with scams (e.g., "won," "money," "click").

ML Model

The model compares these features against millions of known spam and ham (legitimate) emails.

Prediction

The model outputs a high probability of Spam, and the email is diverted from the inbox.

TESTING KNOWLEDGE: IN-CLASS ACTIVITY

Quick questions to verify our understanding of today's core concepts

01

Q1: Which is NOT an NLP application?

(A) Machine Translation, (B) Sentiment Analysis, (C) Image Compression, (D) Spam Detection. Answer: C

02

Q2: Which NLP task is primary to ChatGPT?

(A) Text Generation, (B) Image Segmentation, (C) Object Detection, (D) Circuit Design. Answer: A

03

Q3: Why is NLP computationally difficult?

(A) Human language is ambiguous, (B) Computers dislike English, (C) There are too many alphabets. Answer: A

CONTINUING THE JOURNEY: SUGGESTED READING

Speech and Language Processing

Jurafsky & Martin (3rd ed. draft). The definitive guide to the field.

Natural Language Processing with Python

Bird, Klein & Loper. A practical guide using the NLTK library.

Introduction to Natural Language Processing

Jacob Eisenstein. A mathematically grounded modern textbook.

Foundations of Statistical NLP

Manning & Schütze. Essential for understanding the probabilistic roots of the field.

Deepen your expertise with these foundational NLP texts

SUMMARIZING OUR NLP FOUNDATION

Bridging linguistics and AI to build the next generation of intelligence